

Modeling Ideas: Transit Modeling for Kepler Light Curves

Paul Baines

June 17, 2013

1 Notation & Preliminaries

Disclaimer: These are just preliminary ideas. There are likely to be typos etc. throughout the document. Use at your own risk. ☺

Consider a random process Y defined as a function of time and depending on unknown parameter(s) θ i.e., $Y(t|\theta)$. For simplicity let $\theta = (\omega, \psi, \xi)$ and \mathcal{T} be the interval on which the process is defined. We assume that:

$$\begin{aligned} Y(t|\theta) &= m(t|\omega, \xi) + \epsilon(t|\psi), \quad \forall t \in \mathcal{T} \\ m(t|\omega, \xi) &= [1 - q(t|\omega)] f(t|\xi), \end{aligned}$$

where $\epsilon(t|\psi)$ is a mean-zero random process and $f(t|\omega)$ is the underlying intensity if there are no transits i.e., $q(t|\omega) \equiv 0$ if there are no transits. In the event of a transit we observe a version of $f(t|\xi)$ ‘dampened’ by a multiplicative factor $1 - q(t|\omega)$. There is a trade-off in the complexity we place in f and in ϵ , but for our purposes we assume a mean-zero Gaussian process for ϵ :

$$\epsilon_t \sim \text{GP}(0, \Sigma_\epsilon(\psi)), \quad (1)$$

where $\Sigma_\epsilon(t_i, t_j) = K(|t_i - t_j|; \psi)$ i.e., stationary covariance. Although the form in (??) is fairly rich, in practice we will often assume either independent noise (i.e., $\Sigma(t_i, t_j) = \sigma_\epsilon^2 \delta_{t_i t_j}$) or Ornstein-Uhlenbeck/AR(1)-type noise. Although we only observe a noisy version of $m(t)$, the goal is to ‘separate’ m into two pieces: one representing ‘transit-like variation’, the other representing ‘everything else’. Fortunately we are able to approximate the form of q parametrically. Let t_0 denote the time of the first transit, t_d the transit duration, α the multiplicative reduction in the signal and P the period. Then we assume:

$$q(t|\omega) = \begin{cases} 0 & \text{if } (t \bmod P) \notin [t_0, t_0 + t_d] \\ \alpha & \text{if } (t \bmod P) \in [t_0, t_0 + t_d] \end{cases},$$

where $\omega = \{t_0, t_D, \alpha, P\}$. More sophisticated models for the transit could also be used to allow for smooth transits. The remaining modeling task is to decide upon a structure for f , for which flexible models are required to capture all possible instrumental and astrophysical signals. We decide to model f in the wavelet domain. Let \mathbf{w} denote the DWT of f and let g denote a known thresholding function (hard or soft) applied to the wavelet coefficients then we could model f as:

$$f(\cdot|\mathbf{w}) = \mathcal{W}^{-1} \mathbf{w}, \quad \text{s.t.} \quad g(\mathbf{w}) = c. \quad (2)$$

For example, g could be used to constrain the high-resolution wavelet coefficients to zero. This can be just be viewed as a modeling the DWT-transformed mean as a linear function of a subset of the wavelet coefficients (i.e., or more generally, just a flexible parametric form for f parametrized by the wavelet coefficients). This full model has parameters $\theta = \{\psi, \omega, \mathbf{w}\}$, corresponding to the noise

(co)variance parameters, the transit parameters and the (denoised) wavelet coefficients describing the noiseless and transit-removed signal. The full model is then just:

$$\begin{aligned} Y(t|\theta) &= \text{GP}([1 - q(t|\omega)] f(t|\xi), \Sigma_\epsilon(\psi)), \\ L(\theta|y) &= \mathcal{N}(y; [1 - q(\cdot|\omega)] f(\cdot|\xi), \Sigma_\epsilon(\cdot; \psi)). \end{aligned} \quad (3)$$

In a likelihood framework direct maximization of (??) with respect to θ is likely to be very difficult (and given the high dimension of \mathbf{w} , probably not advisable). However, given there may be prior information available about certain parameters we will instead take a more flexible approach that actually leads to simpler computation. Thus, instead of applying hard constraints on the wavelet coefficients we instead propose to model:

$$\mathcal{W}\mathbf{f} = \mathbf{w}, \quad \mathbf{w}|d, \Sigma_w \sim N(d, \Sigma_w), \quad \mathbf{d}|\Sigma_d \sim N(0, \Sigma_d).$$

where \mathcal{W} represents the DWT in matrix form. Thus:

$$\mathbf{d}|\mathbf{w}, \Sigma_w, \Sigma_d \sim N\left([\Sigma_d^{-1} + \Sigma_w^{-1}]^{-1} [\Sigma_w^{-1} \mathbf{w}], [\Sigma_d^{-1} + \Sigma_w^{-1}]^{-1}\right).$$

This illustrates the shrinkage of the wavelet coefficients toward zero by an amount controlled by Σ_w . Typically Σ_w would be specified to ensure a large amount of shrinkage for high-resolution coefficients, and minimal shrinkage for low-resolution coefficients (and would be diagonal). In practice as well, variances can be pooled across resolution levels to yield:

$$w_{km} \sim N(d_{km}, \sigma_{w,m}^2), \quad d_{km} \sim N(0, \sigma_{d,m}^2), \quad \sigma_{w,m}^2 \sim \text{Inv-}\chi^2(\nu_0, s^2).$$

Thus:

$$d_{km}|w_{km}, \sigma_{w,m}^2, \sigma_{d,m}^2 \sim N\left(\frac{(\sigma_{w,m}^2)^{-1} w_{km}}{(\sigma_{d,m}^2)^{-1} + (\sigma_{w,m}^2)^{-1}}, \frac{1}{(\sigma_{d,m}^2)^{-1} + (\sigma_{w,m}^2)^{-1}}\right).$$

For levels with little or no shrinkage we obtain:

$$d_{km}|w_{km}, \sigma_{w,m}^2, (\sigma_{d,m}^{-2} \approx 0) \stackrel{\text{approx.}}{\sim} N(w_{km}, \sigma_{w,m}^2),$$

and for levels with a high degree of shrinkage, the posterior for d_{km} is concentrated around zero. For the time being, let's just consider Σ_d (and possibly Σ_w as well) to be specified by the analyst. Note that \mathbf{f} and \mathbf{w} are one-to-one so that:

This version of the full model either has parameters $\theta = \{\psi, \omega\}$ or $\theta = \{\psi, \omega, \Sigma_w^2\}$ depending upon how the wavelet coefficients are modeled. The likelihood is given by:

$$\begin{aligned} p(y, f, d|\theta) &= p(y|f, \omega, \psi) p(f|d, \Sigma_w^2) p(d), \\ \Rightarrow p(y|\theta) &= \int p(y, f, d|\theta) df dd, \end{aligned} \quad (4) \quad (5)$$

again, the full Bayesian version simply requires a prior on all components of θ and may be preferable.

$$\begin{aligned} p(f, d, \theta|y) &= p(y|f, \omega, \psi) p(f|d, \Sigma_w^2) p(d) p(\psi, \omega, \Sigma_w^2), \\ \Rightarrow p(\theta|y) &= \int p(f, d, \theta|y) df dd, \end{aligned} \quad (6) \quad (7)$$

The form of (??) lends itself to a natural Gibbs sampler of the form:

$$[\psi|f, \omega], \quad [\omega|\psi, f], \quad [f|d, \omega, \psi, \Sigma_w^2], \quad [d|f, \Sigma_w^2], \quad [\Sigma_w^2|d, f],$$

where the conditioning on y throughout is dropped for brevity. Let us briefly describe each of the Gibbs updates:

- Sampling from $\psi|f, \omega$: This just amounts to sampling from the posterior distribution of the covariance parameters of a Gaussian process. If the covariance is diagonal then, with a conjugate prior, ψ can be sampled exactly (Inverse- χ^2). If ψ includes covariance parameters then this can be done using e.g., MH.
- Sampling from $\omega|\psi, f$: This will not be in closed form and other methods (e.g., MH) will be needed. Intuitively, varying ω given f varies the value of $m(t)$, so values of ω that yield m to be an appropriate mean for y (given ψ) will be favored. Likely to be the most challenging step to effectively explore the posterior.
- Sampling from $f|d, \omega, \psi, \Sigma_w^2$: Since f is sandwiched within the multilevel model, the ‘prior’ on f is multivariate normal, and f appears linearly in the ‘likelihood’, this will also be multivariate normal (albeit with some linear algebra and an inverse DWT required to compute the mean and variance).
- Sampling from $d|f, \Sigma_w^2$: This is just multivariate normal. The current value of f is transformed to the wavelet domain, and then the mean of d is given by a shrunk version of the wavelet coefficients corresponding to the shrinkage induced by the relative weights of Σ_w^2 and Σ_d^2 .
- Sampling from $\Sigma_w^2|d, f$: This just amounts to sampling from the posterior distribution of the covariance parameters of a multivariate normal. If the covariance is diagonal (which it will be) or fully unstructured then, with a conjugate prior this can be sampled exactly (Inverse- χ^2). For parametrized non-diagonal matrices it can be sampled using e.g., MH.

To make things more concrete, consider the special case with n evenly-spaced observations (no missing data),

$$\Sigma_\epsilon(\psi) = \sigma^2 I, \quad \Sigma_d = \text{diag}(\sigma_{d,1}^2, \dots, \sigma_{d,1}^2, \sigma_{d,2}^2, \dots, \sigma_{d,J}^2), \\ \Sigma_w = \text{diag}(\sigma_{w,1}^2, \dots, \sigma_{w,1}^2, \sigma_{w,2}^2, \dots, \sigma_{w,J}^2),$$

where the elements of Σ_d are specified constants. Full details of the sampling procedure for this special case are given below:

- Sampling from $\sigma^2|f, \omega$: With a conjugate prior $\sigma^2 \sim \text{Inv-}\chi^2(\nu_0, s_0^2)$ on σ^2 we obtain the posterior:

$$\sigma^2|f, \omega \sim \text{Inv-}\chi^2(\nu_n, s_n^2),$$

where

$$\nu_n = \nu_0 + n, \quad s_n^2 = \frac{\nu_0 s_0^2 + (y - f)^T (y - f)}{\nu_n}.$$

- Sampling from $\omega|\psi, f$: This will not be in closed form and other methods (e.g., MH) will be needed. Intuitively, varying ω given f varies the value of $m(t)$, so values of ω that yield m to be an appropriate mean for y (given ψ) will be favored.
- Sampling from $f|d, \omega, \psi, \Sigma_w^2$: Since,

$$f|d, \Sigma_w \sim N(\mathcal{W}^{-1}d, \mathcal{W}^{-1}\Sigma_w\mathcal{W}^{-T}), \quad y|f \sim N((I - Q)f, \Sigma_\epsilon)$$

Therefore:

$$f|y, d, \Sigma_w \sim N(\mu_f, \Sigma_f), \\ \Sigma_f^{-1} = \mathcal{W}^T \Sigma_w^{-1} \mathcal{W} + (I - Q) \Sigma_\epsilon^{-1} (I - Q), \\ \mu_f = \Sigma_f [\mathcal{W}^T \Sigma_w^{-1} d + (I - Q) \Sigma_\epsilon^{-1} y].$$

Note that if $\Sigma_w = \sigma_w^2 I$ then Σ_f is diagonal and $\mathcal{W}^T d$ is the IDWT of d .

- Sampling from $d|f, \Sigma_w^2$: Let $w = \mathcal{W}f$, then:

$$\mathbf{d}|\mathbf{w}, \Sigma_w, \Sigma_d \sim N\left([\Sigma_d^{-1} + \Sigma_w^{-1}]^{-1} [\Sigma_w^{-1} \mathbf{w}], [\Sigma_d^{-1} + \Sigma_w^{-1}]^{-1}\right).$$

Note that both Σ_w and Σ_d are diagonal, so there are no matrix inversions in this step.

- Sampling from $\Sigma_w^2|d, f$: With independent conjugate priors on each $\sigma_{w,j}^2 \sim \text{Inv-}\chi^2(\nu_{w,j,0}, s_{w,j,0}^2)$ on each $\sigma_{w,j}^2$ for $j = 1, \dots, J$ we obtain:

$$\sigma_{w,j}^2|d, f \sim \text{Inv-}\chi^2(\nu_{w,j,n}, s_{w,j,n}^2),$$

where

$$\nu_{w,j,n} = \nu_{w,0} + n_j, \quad s^2 = \frac{\nu_{w,0} s_{w,j,0}^2 + (w_j - d_j)^T (w_j - d_j)}{\nu_{w,j,n}}.$$

Here n_j is the number of coefficients in level j of the wavelet decomposition and w_j and d_j are the subcomponents of w and d corresponding to the wavelet coefficients at level j .

2 Missing Data

Handling missing data in the time domain. Straightforward: write up.

3 Searching for Transits: Finding Candidate Period/Phase Configurations

4 Using Metadata to Improve Performance

Any ideas?

5 Model Comparison

- Ideas to ‘test’ for presence of transits
- Stellar variability models.
- Spacecraft event models (e.g., sudden pixel dropout)